

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 import datetime as dt
6
7 import sklearn
8 from sklearn.preprocessing import StandardScaler
9 from sklearn.cluster import KMeans
10 from sklearn.metrics import silhouette_score
```

```
1 from google.colab import drive
2 drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
1 df=pd.read_csv("/content/drive/MyDrive/Colab Notebooks/OnlineRetail.csv",encoding="ISO-8859-1")
2 df
```

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Count
0	536365	85123A	WHITE HANGING HEART T- LIGHT HOLDER	6	01-12-2010 08:26	2.55	17850.0	Unit Kingdo
1	536365	71053	WHITE METAL LANTERN	6	01-12-2010 08:26	3.39	17850.0	Unit Kingdo
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	01-12-2010 08:26	2.75	17850.0	Unit Kingdo
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	01-12-2010 08:26	3.39	17850.0	Unit Kingdo
RED WOOLLY								

```
1 df.shape
```

(49445, 8)

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 49445 entries, 0 to 49444
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  -
0   InvoiceNo    49445 non-null  object
1   StockCode   49445 non-null  object
2   Description  49303 non-null  object
3   Quantity    49445 non-null  int64
4   InvoiceDate  49445 non-null  object
5   UnitPrice   49445 non-null  float64
6   CustomerID  31563 non-null  float64
7   Country     49444 non-null  object
dtypes: float64(2), int64(1), object(5)
memory usage: 3.0+ MB
```

```
1 df.describe()
```

	Quantity	UnitPrice	CustomerID
count	49445.000000	49445.000000	31563.000000
mean	8.315239	6.458260	15423.016823
std	55.832343	167.436019	1753.298111
min	-9360.000000	0.000000	12347.000000
25%	1.000000	1.250000	14051.000000
50%	2.000000	2.510000	15464.000000
75%	8.000000	4.250000	17041.000000
max	2880.000000	16888.020000	18283.000000



```
1 df.isnull().sum()
```

```
InvoiceNo      0
StockCode      0
Description    142
Quantity       0
InvoiceDate    0
UnitPrice      0
CustomerID    17882
Country        1
dtype: int64
```

```
1 df=df.dropna()
2 df.isnull().sum()
```

```
InvoiceNo      0
StockCode      0
Description     0
Quantity       0
InvoiceDate    0
UnitPrice      0
CustomerID     0
Country        0
dtype: int64
```

```
1 df.shape
```

```
(31563, 8)
```

```
1 df['CustomerID']=df['CustomerID'].astype(str)
```

```
<ipython-input-17-d4b6fccbd77f>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df['CustomerID']=df['CustomerID'].astype(str)

```
1 df['Amount']=df['UnitPrice']*df['Quantity']
2 fm_m=df.groupby('CustomerID')['Amount'].sum()
3 fm_m=fm_m.reset_index()
4 fm_m.columns=['CustomerID','Spend_Amount']
5 fm_m.head()
```

```
<ipython-input-18-523d3fc75744>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df['Amount']=df['UnitPrice']*df['Quantity']

	CustomerID	Spend_Amount
--	------------	--------------

0	12347.0	711.79
---	---------	--------

1	12348.0	892.80
---	---------	--------

2	12370.0	1868.02
---	---------	---------

3	12377.0	1001.52
---	---------	---------

4	12383.0	600.72
---	---------	--------



```
1 fm_f=df.groupby('CustomerID')['InvoiceNo'].count()
2 fm_f=fm_f.reset_index()
3 fm_f.columns=['CustomerID','Frequency']
4 fm_f.head()
```

	CustomerID	Frequency
--	------------	-----------

0	12347.0	31
---	---------	----

1	12348.0	17
---	---------	----

2	12370.0	91
---	---------	----

3	12377.0	43
---	---------	----

4	12383.0	37
---	---------	----




```
1 fm=pd.merge(fm_m,fm_f,on='CustomerID',how='inner')
2 fm.head()
```

	CustomerID	Spend_Amount	Frequency
0	12347.0	711.79	31
1	12348.0	892.80	17
2	12370.0	1868.02	91
3	12377.0	1004.50	42



```
1 attributes=['Spend_Amount','Frequency']
2 plt.rcParams['figure.figsize']=[7,7]
3 sns.boxplot(data=fm[attributes],orient="X",palette="Set2", whis=1.5,saturation=1,width=0.7)
4 plt.title("Outliers Variable Distribution",fontsize=14,fontweight='bold')
5 plt.ylabel("Range",fontweight='bold')
6 plt.xlabel("Attributes",fontweight='bold')
7
```

Text(0.5, 0, 'Attributes')



```
1 fm.shape
```

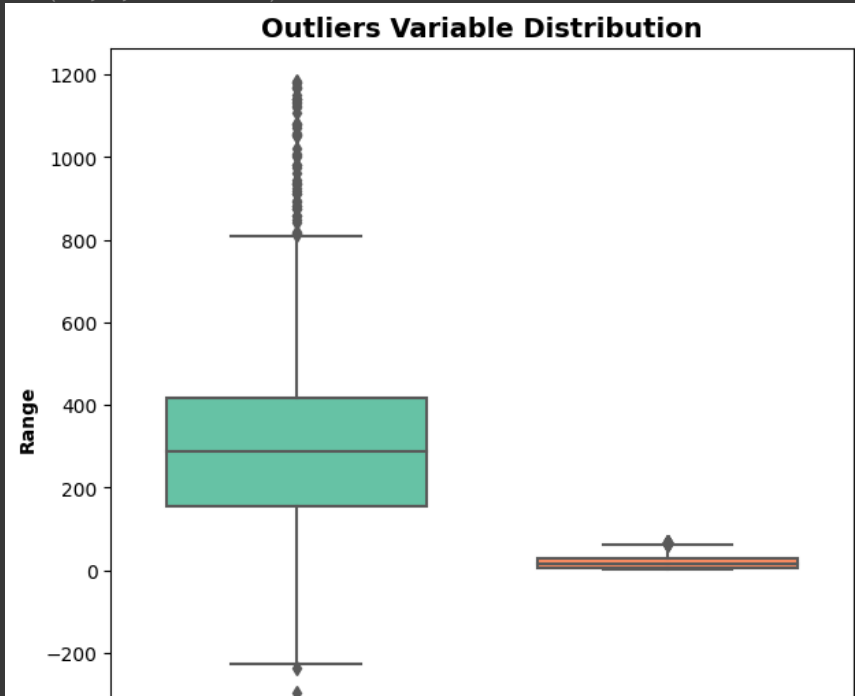
(1027, 3)

ChatGPT

```
1 Q1=fm.Spend_Amount.quantile(0.25)
2 Q3=fm.Spend_Amount.quantile(0.75)
3 IQR=Q3-Q1
4 fm=fm[(fm.Spend_Amount>=Q1-1.5*IQR)& (fm.Spend_Amount<=Q3+1.5*IQR)]
5
6 Q1=fm.Frequency.quantile(0.25)
7 Q3=fm.Frequency.quantile(0.75)
8 IQR=Q3-Q1
9 fm=fm[(fm.Frequency>=Q1-1.5*IQR)& (fm.Frequency<=Q3+1.5*IQR)]
10
11
```

```
1 attributes=['Spend_Amount','Frequency']
2 plt.rcParams['figure.figsize']=[7,7]
3 sns.boxplot(data=fm[attributes],orient="X",palette="Set2", whis=1.5,saturation=1,width=0.7)
4 plt.title("Outliers Variable Distribution",fontsize=14,fontweight='bold')
5 plt.ylabel("Range",fontweight='bold')
6 plt.xlabel("Attributes",fontweight='bold')
7
```

Text(0.5, 0, 'Attributes')



```
1 fm.shape
```

```
(858, 3)
```

Spend_Amount Frequency

```
1 fm.describe()
```

	Spend_Amount	Frequency
count	858.000000	858.000000
mean	323.150501	20.074592
std	258.816060	16.728919
min	-442.500000	1.000000
25%	155.687500	7.000000
50%	288.600000	16.000000
75%	417.540000	29.000000
max	1183.090000	70.000000

```
1 fm_df=fm[['Spend_Amount','Frequency']]
2 scaler=StandardScaler()
3 fm_df_scaled=scaler.fit_transform(fm_df)
4 fm_df_scaled.shape
```

```
(858, 2)
```

```
1 fm_df_scaled
```

```
array([[ 1.50248086,  0.65346602],
       [ 2.20226579, -0.18389624],
       [ 2.62257746,  1.3712051 ],
       ...,
       [-0.35741377, -0.54276578],
       [-0.7123127 , -0.72220055],
       [-0.83003245,  2.20856736]])
```

```
1 fm_df_scaled=pd.DataFrame(fm_df_scaled)
2 fm_df_scaled.columns=['Amount','Frequency']
3 fm_df_scaled.head()
```

	Amount	Frequency
0	1.502481	0.653466
1	2.202266	-0.183896
2	2.622577	1.371205
3	1.073084	1.012336
4	0.304446	-0.602577


```
1 fm_df_scaled.describe()
```

	Amount	Frequency
count	8.580000e+02	8.580000e+02
mean	1.656277e-17	-7.039176e-17
std	1.000583e+00	1.000583e+00
min	-2.960006e+00	-1.140882e+00
25%	-6.474122e-01	-7.820121e-01
50%	-1.335723e-01	-2.437078e-01
75%	3.649099e-01	5.338428e-01
max	3.324527e+00	2.986118e+00

```
1 track_inertia={}
2 silhouette={}
3 for k in range(2,11):
4     kmeans=KMeans(n_clusters=k,init='k-means++',random_state=0)
5     kmeans.fit(fm_df_scaled)
6     track_inertia[k]=kmeans.inertia_
7     silhouette[k]=silhouette_score(fm_df_scaled,kmeans.labels_)
```


```
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'a
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'a
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'a
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'a
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'a
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'a
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'a
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'a
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'a
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'a
warnings.warn(
```

```
1 sns.pointplot(x=list(track_inertia.keys()),y=list(track_inertia.values()))
```



```
1 sns.pointplot(x=list(silhouette.keys()),y=list(silhouette.values()))
```

<Axes: >



ChatGPT

```
1 kmeans=KMeans(n_clusters=3,init='k-means++',max_iter=300,n_init=10,random_state=0)
2 kmeans.fit(fm_df_scaled)
3 print("WCSS for k: %d and silhouette score: %f" %(kmeans.inertia_,silhouette_score(fm_df_scaled,kmeans.labels_)))
4
```

WCSS for k: 603 and silhouette score: 0.484208


```
1 kmeans.cluster_centers_
```

```
array([[ 1.94709735,  0.71176975],
       [-0.43576192, -0.56482396],
       [ 0.08037223,  1.33634349]])
```

```
1 fm['Cluster_Id']=kmeans.labels_
2 fm.head()
```


	CustomerID	Spend_Amount	Frequency	Cluster_Id
0	12347.0	711.79	31	0
1	12348.0	892.80	17	0
3	12377.0	1001.52	43	0
4	12383.0	600.72	37	0
5	12386.0	401.90	10	1

```
1 plt.figure(figsize=(7,7))
2 plt.scatter(fm_df_scaled["Amount"],fm_df_scaled["Frequency"],c=kmeans.labels_,cmap='viridis')
3 plt.scatter(kmeans.cluster_centers_[ :,0],kmeans.cluster_centers_[ :,1],c='blue',s=200,marker='x')
4 plt.xlabel('Amount')
5 plt.ylabel('Frequency')
6 plt.title('Scatter Plot of Amount vs Frequency')
7 plt.show()
```



```
1 sns.relplot(x='Spend_Amount',y='Frequency', data=fm,hue=fm['Cluster_Id'],height=5)
```

<seaborn.axisgrid.FacetGrid at 0x783df54f3a60>



```
1 new_data_point=[390.5,23]
2 new_data_point_scaled=scaler.transform([new_data_point])
3 predicted_cluster=kmeans.predict(new_data_point_scaled)
4 print(f"The new data point belongs to cluster {predicted_cluster[0]}")
```

The new data point belongs to cluster 1

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but KMeans was fitted with feature names
warnings.warn(
)